

# Chapter 2, Section 2.1 to 2.3: Architecture and Characteristics of a RTOS kernel ( $\mu$ C/OS-II)

*1. Introduction*

*2. Interrupt Management*

*3. Task Management*

**4. Time Management**

**5. Event Management**

**6. Memory Management**


**7. Porting  $\mu$ COS-II**

# 1. Introduction

- $\mu\text{C/OS-II}$  is one of the Micrium products (<http://www.micrium.com/>)
- $\mu\text{C/OS-II}$  is a real-time multitasking kernel, scalable and easily adaptable to different processors.
- It can be programmed in a ROM or a Flash memory.
- Also, latencies of interrupt and switching context are deterministic and low ( $O(\mu\text{s})$ ).

# 1. Introduction

- $\mu\text{C}/\text{OS-II}$  is second after VxWorks (percentage at the level of the RTOS market ).
- $\mu\text{C}/\text{OS-II}$  offers one the most popular books for the embedded systems topic.



# 1. Introduction

- **$\mu$ C/OS-II source and object code can be used by accredited Colleges and Universities without requiring a license, as long as there is no commercial application involved. In other words, no licensing is required if  $\mu$ C/OS-II and/or  $\mu$ C/TCP-IP are used for educational use.**
- **You need to obtain an 'Executable Distribution License' to embed  $\mu$ C/OS-II in a product that is sold with the intent to make a profit or if the product is not used for education or 'peaceful' research.**

# 1. Introduction

- **See also:**
  - **the list of products:**  
<http://www.micrium.com/products/uc-products.html>
  - **the partial list of client:**  
<http://www.micrium.com/customers/list.html>
  - **the list of porting:**  
<http://www.micrium.com/products/rtos/kernel/ports.html>

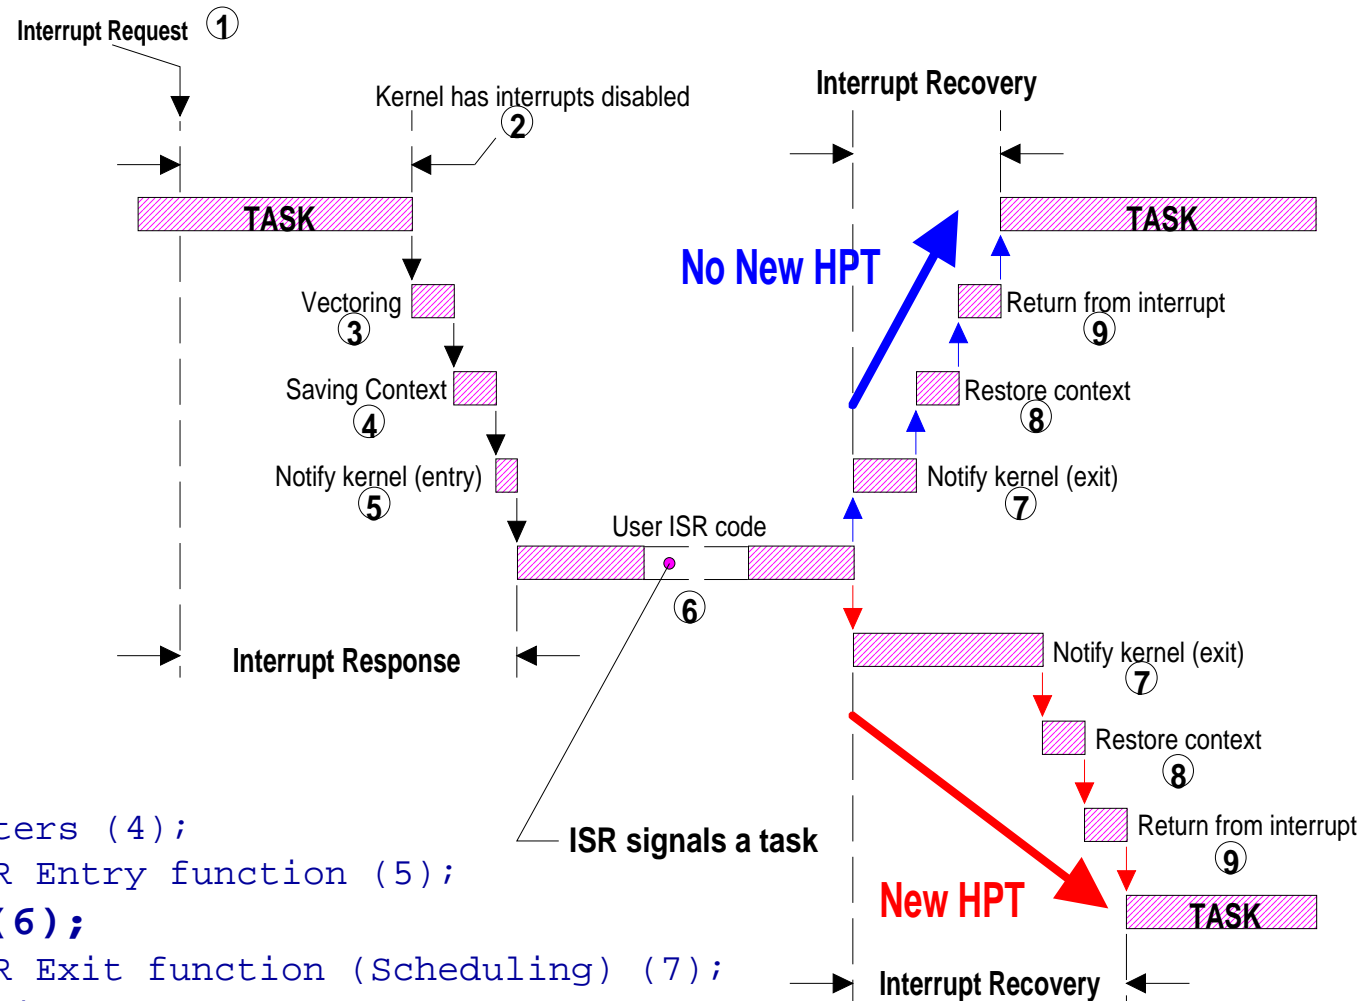
## 2. Interrupts management

- **Your Interrupt Service Routine (ISR)...**
  - **Saves all the CPU registers;**
  - Calls OSIntEnter() or increment OSIntNesting directly;**
  - If (OSIntNesting == 1)**
    - OSTCBCur -> OSTCBStkPtr = SP; //details later**
  - Clears interrupting device;**
  - Re-enables interrupts (optional)**
  - Executes user code to service interrupt;**
  - Calls OSIntExit();**
  - Restores all CPU registers**
  - Executes a return from interrupt instruction;**

## 2 Interruptions management

- **OSIntEnter() :**
  - Disables interrupts to gain exclusive access to the global variable OSIntNesting (current number of interrupt nested)
  - Increments OSIntNesting
  - Enables interrupts
- **OSIntExit() :**
  - Disables interrupts
  - Decrements OSIntNesting
  - If OSIntNesting == 0, determines HPT (Highest Priority Task) as the next task to be executed
  - Performs a context switch to resume HPT
  - Enables interrupts;

Time →



ISR (3):

Save CPU Registers (4);

Call Kernel ISR Entry function (5);

**Process ISR (6);**

Call Kernel ISR Exit function (Scheduling) (7);

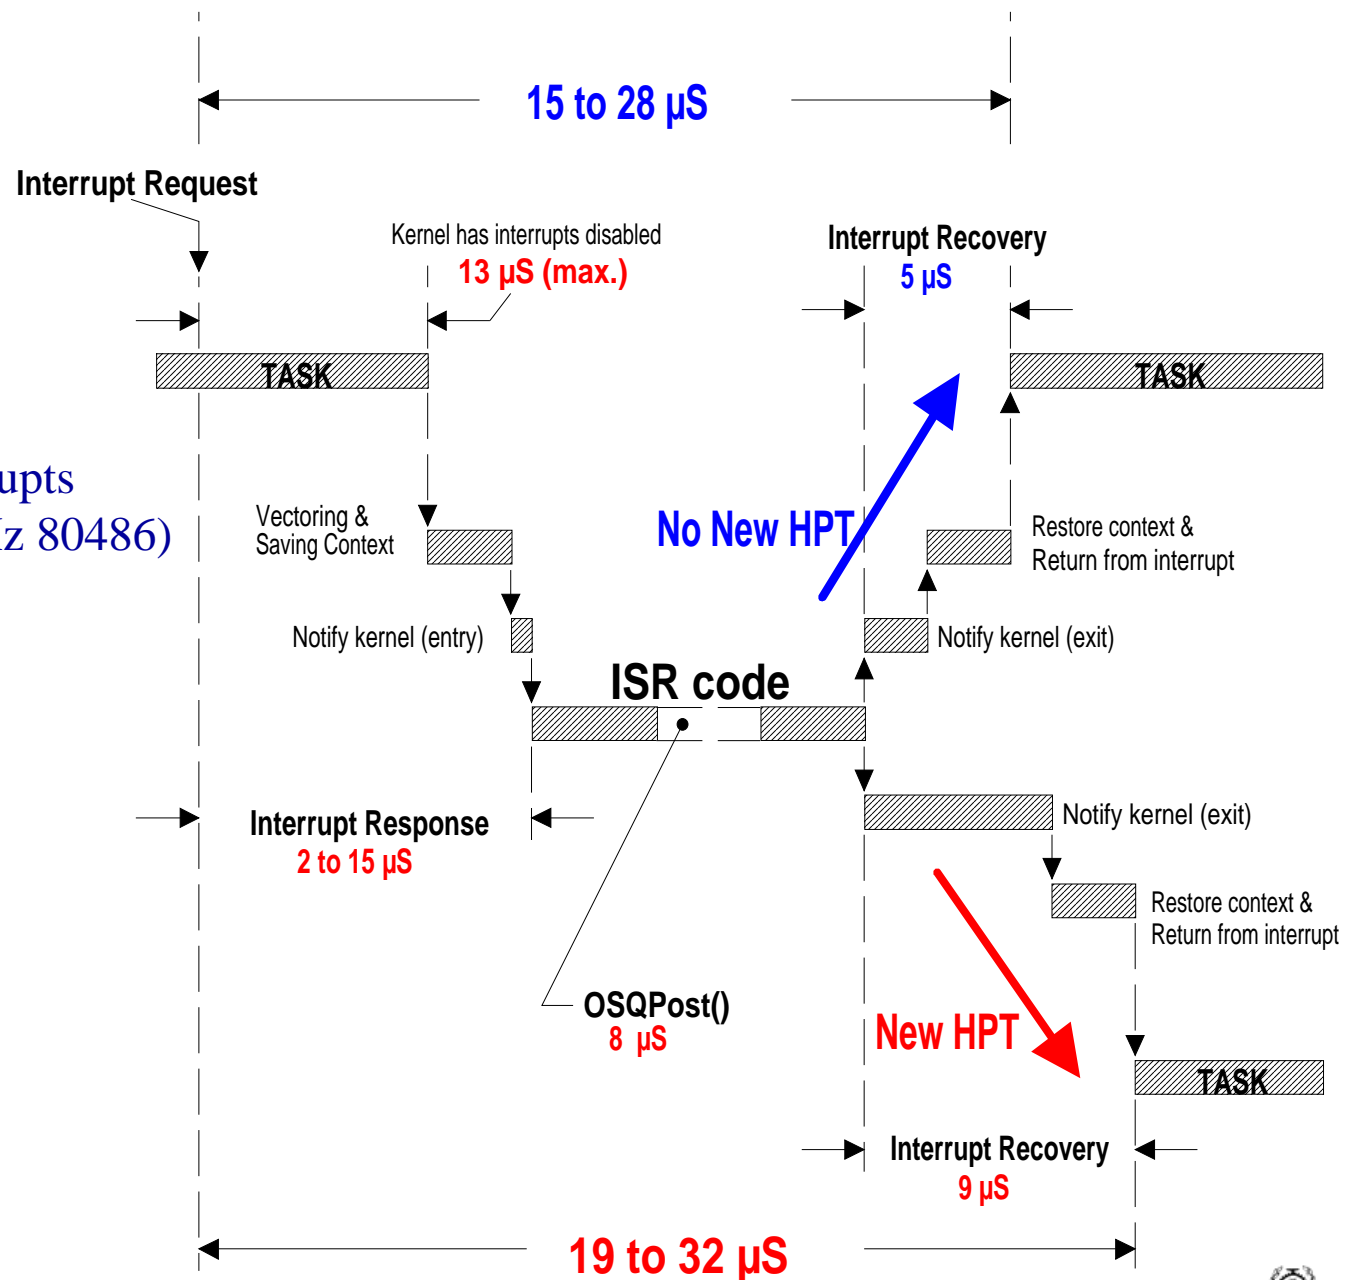
Restore CPU Registers (8);

Return from Interrupt (9);

Real-Time Kernels / Copyright 2001, Jean J. Labrosse



## Servicing Interrupts ( $\mu$ C/OS-II on 66 MHz 80486)



# 3. Task Management

## Overview

- 3.1 A task and its data structure
- 3.2 Task states
- 3.3 Creating a task
- 3.4 Deleting a task
- 3.5 Suspending/resuming a task
- 3.6 Delaying a task
- 3.7 Changing the priority of a task
- 3.8 Tasks scheduling
- 3.9 System tasks

## 3.1 A task and its data structure


- **A  $\mu$ C/OS-II task is built with an infinite loop**
  - i.e. *for(;;)* or *while(1)*
  - Similar to C subroutines
  - The return type must always be declared *void*
  - The task can delete itself upon completion

## 3.1 A task and its data structure

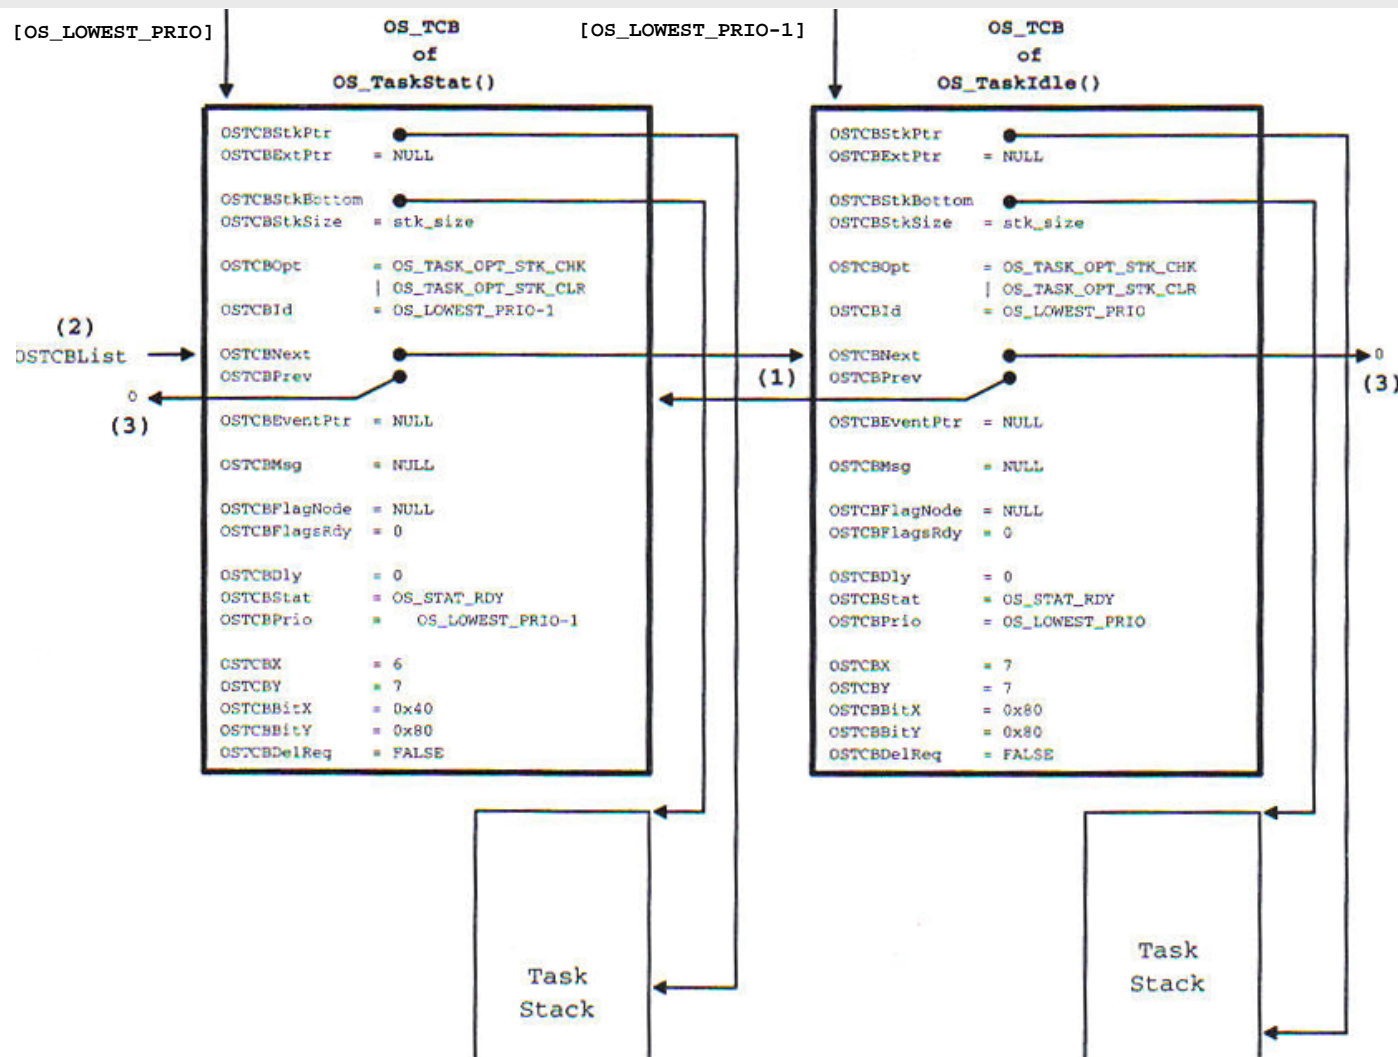
- A task is composed of a TCB (Task Control Block) and a stack
- The TCB contains the task parameters (see next slide)
- The stack contains:
  - functions called by your task;
  - local variables that will be allocated by all functions called by your task,
  - nested interrupts
  - CPU register saved during the context switch

# 3.1 A task and its data structure

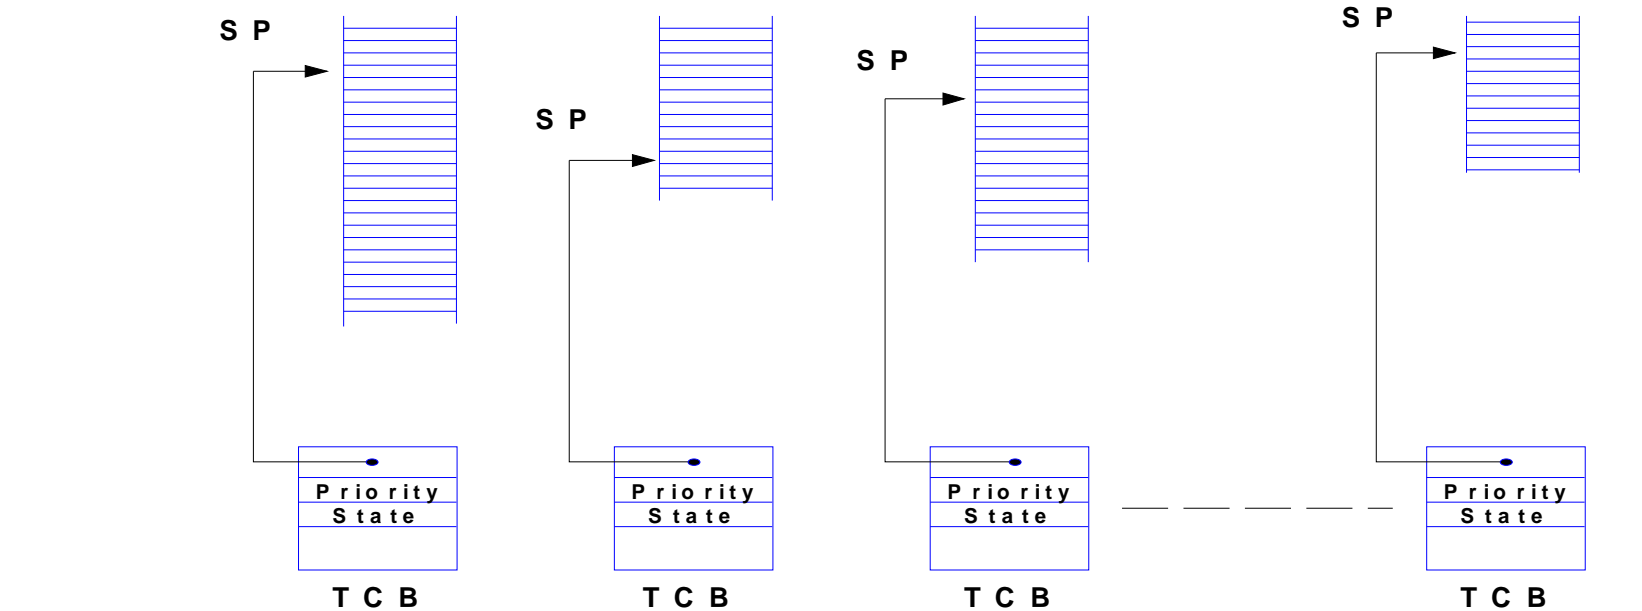
*Variables and data structures after calling OSInit().*



## 3.1 A task and its data structure



# 3.1 A task and its data structure



MEMORY (RAM)

CPU

80x86 CPU  
(Real Mode)

A X
B X
C X
D X

S I
D I
B P

D S
E S
C S
S S

P S W
I P
S P

16 Bit

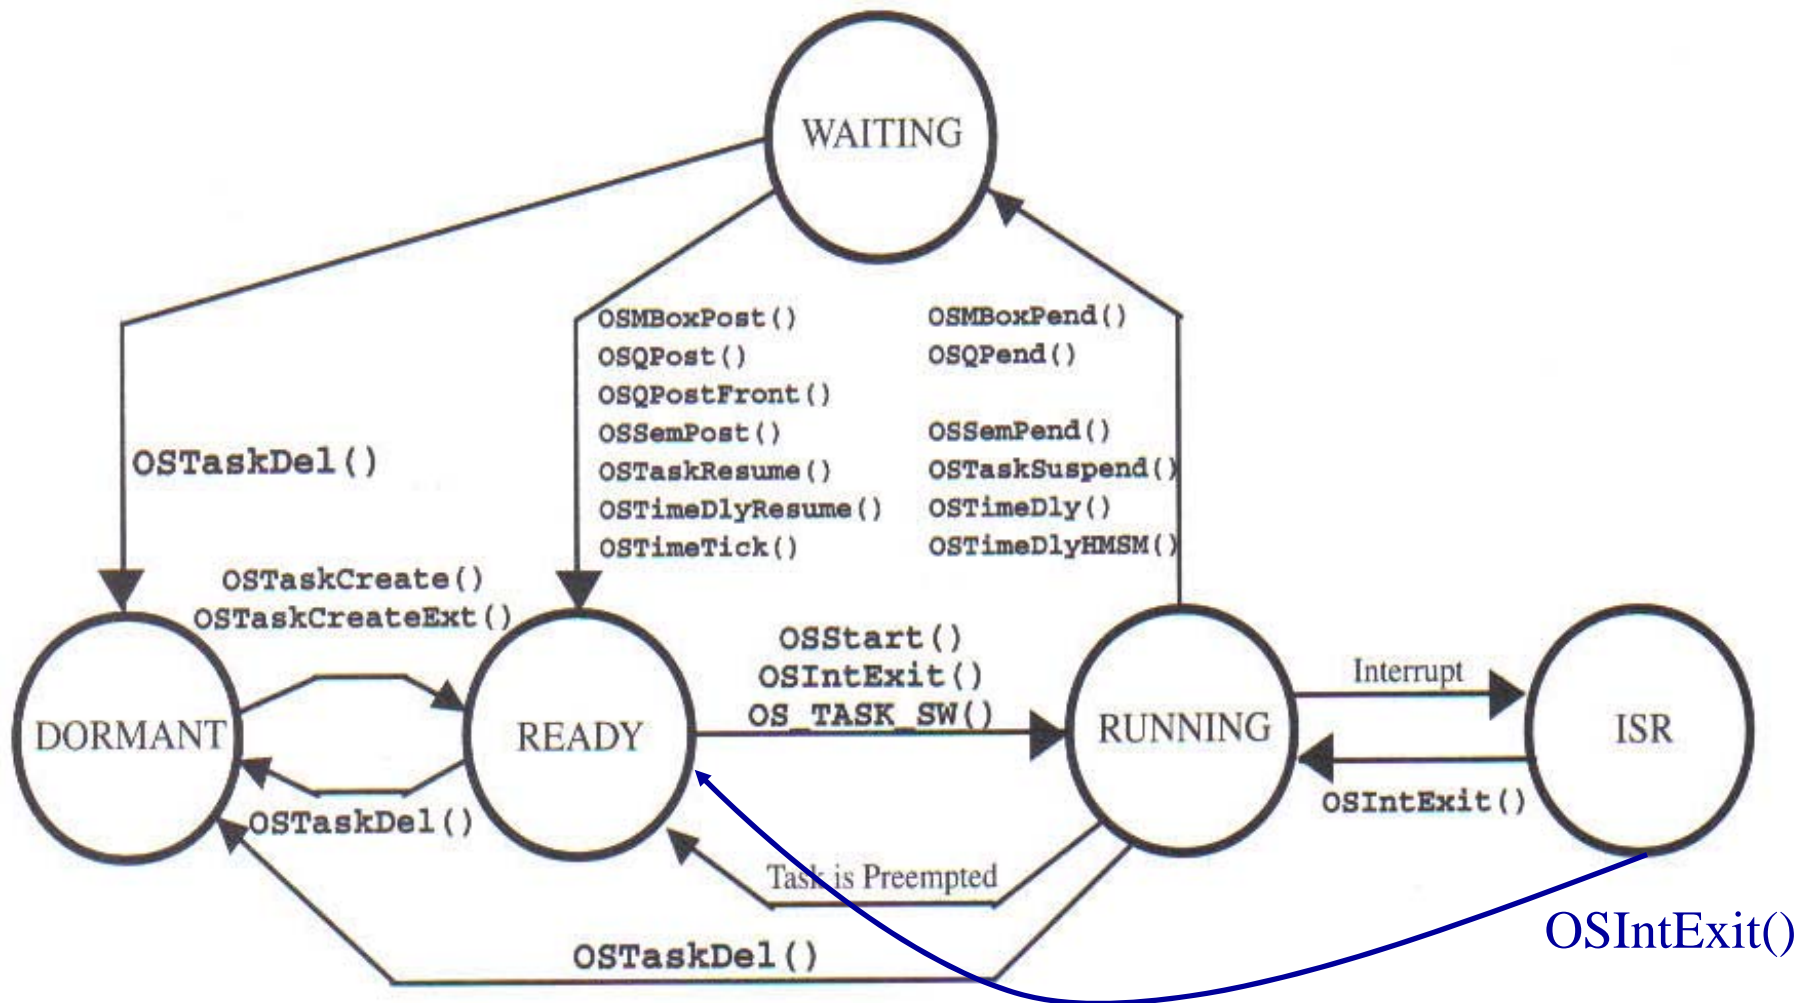
Real-Time Kernels / Copyright 2001, Jean J. Labrosse

Chap2, Sections 2.1 à 2.3, page 15

Copyright© 2011 G. Bois, M. De Nanclas, L. Filion



## 3.2 Task States



Real-Time Kernels / Copyright 2001, Jean J. Labrosse




## 3.3 Creating a task

- **During the creation of a task, different operations are executed by the kernel:**
  - RAM memory allocation in the *System Memory Pool* for the stack and TCB.
  - Stack Initialization (e.g. push the first address of task code and its initial arguments, and also initialize the stack using symbols to facilitate stack checking)
  - TCB Initialization

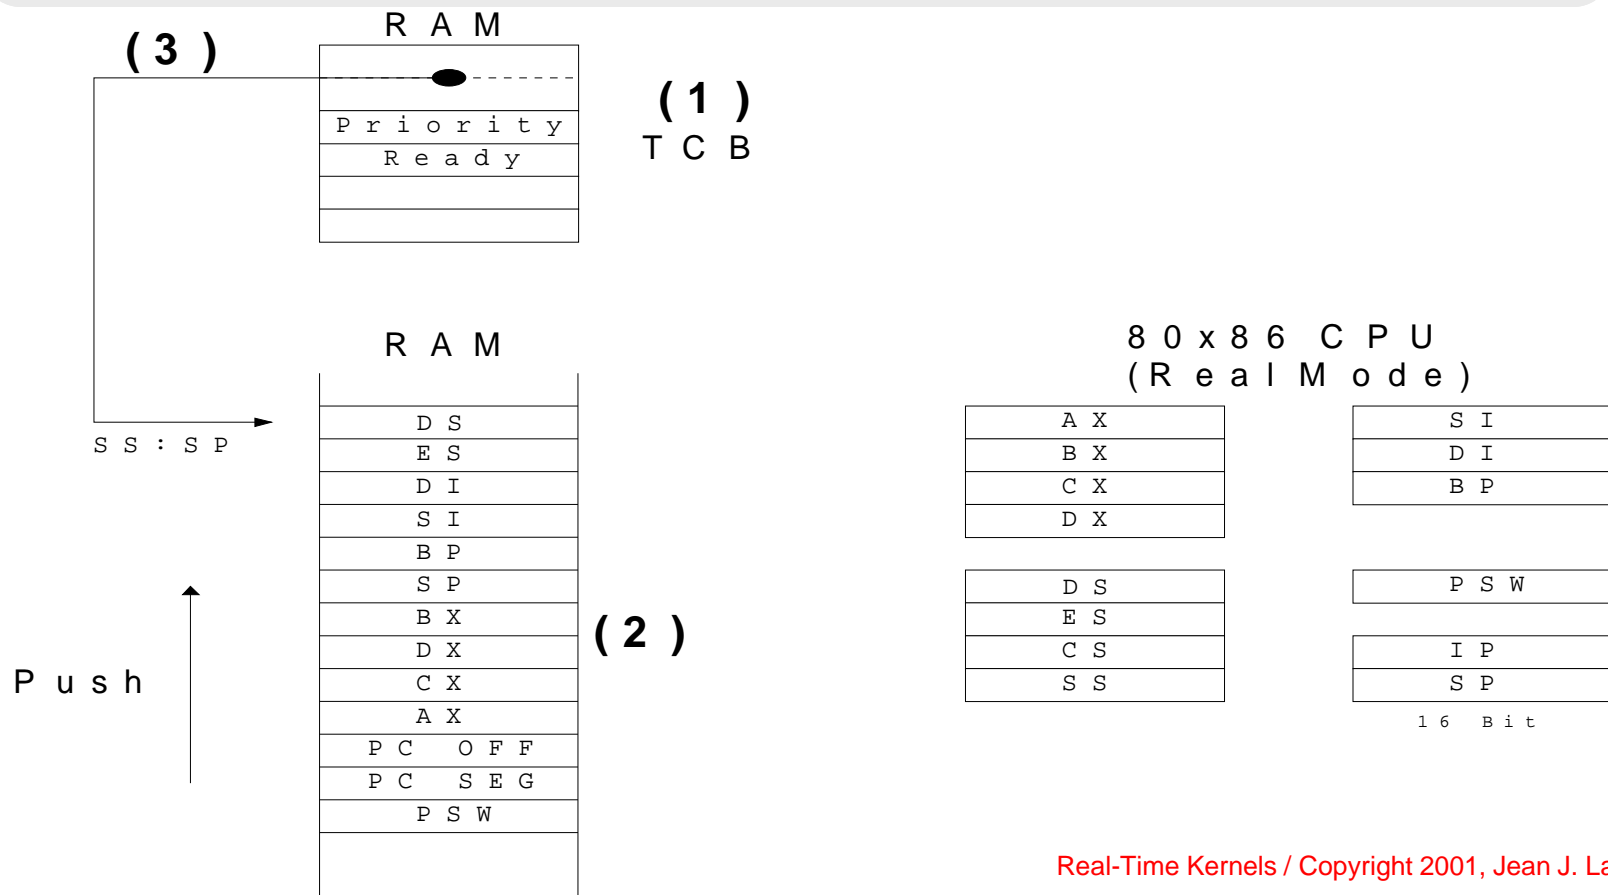
## 3.3 Creating a task

### *Stack checking.*



Real-Time Kernels / Copyright 2001, Jean J. Labrosse

## 3.3 Creating a task



Real-Time Kernels / Copyright 2001, Jean J. Labrosse

- (1) Allocate TCB & Init.
- (2) Init. Task Stack frame
- (3) Save Top-Of-Stack (TOS) in TCB

## 3.3 Creating a task

- **Different ways to create a task:**
- **Examples:**
  - VxWorks: taskCreate()
  - $\mu$ C/OS-II: OSTaskCreateExt() // 9 arguments

## 3.3 Creating a task

```
INT8U OSTaskCreateExt(    void (*task) (void *pd),
                          void *pdata,
                          OS_STK *ptos,
                          INT8U prio,
                          INT16U id,
                          OS_STK *pbos,
                          INT32U stk-size,
                          void *pext,
                          INT16U opt);
```

- task = pointer to the task code
- pdata = pointer to an argument that is passed to your task when on the first execution
- ptos = pointer to the top of the stack
- prio = task priority
- id = unique id of the task (for future expansion)

## 3.3 Creating a task

### *Arguments of OSTaskCreateExt() (cnt'd):*

- ppos = pointer to the task's bottom-of-stack (used to perform stack checking)
- stk\_size = size of the stack. For example, if an entry is 4-byte wide, then a stack size of 1,000 means that the stack has 4,000 bytes.
- pext = pointer to a user-supplied data area that can be used to extend the OS\_TCB of the task (e.g. floating-point registers)
- opt = specific options of OSTaskCreateExt() (see the reference manual)

## 3.3 Creating a task

### *Creation summary*

```
OSTaskCreate (parameters)  
{  
    if (a TCB is available)  
        Initialize the TCB;  
    else  
        Return an error code;  
    Initialize the task's stack;  
    Increment the number of tasks counter;  
    Make task ready-to-run;  
    if (OS is running)  
        Call the scheduler;  
}
```

## 3.3 Creating a task

### *About the size of the stack*

- **The size of the stack needed by the task is application-specific. You must account:**
  - for nesting of all functions called by your task,
  - the number of local variables that will be allocated by all functions called by your task,
  - the stack requirements for all the nested interrupts,
  - the number of all the CPU registers saved during the context switch,
- **A good practice is to allocate twice of the estimation and to execute many times for critical situations. Then, see if you can decrease.**



## 3.3 Creating a task

### *Minimizing the latency*

- **Creating a task increases the execution time and can be a non-deterministic operation**
  - A good practice is to create all the tasks required during the initialization (special functions) and suspend those not immediately required. They will be resumed later.
  - Also if possible (e.g. after debugging) turn the option to determine how much task space is used off (*opt* parameter in *OSTaskCreateExt*).

## 3.3 Creating a task

### *Information about a task*

- Each RTOS provides system functions to retrieve information about a task:
  - name, ID, task code, state, priority, delay.
  - information about the stack
  - options of the task
  - contents of the general registers
  - contents of the floating point registers (FPU)
- *OSTaskQuery()* performs such a function with  $\mu$ C/OS-II.

## 3.4 Deleting a task

- Sometimes, it is necessary to delete a task
- Deleting means that the task is returned to the dormant state
- It is performed in two steps:
  - Remove the task from the ready list and from all the waiting list (e.g. OSReadyTbl and OSEventTbl).
  - Deallocate the TCB and the stack. Deallocated memory can be reused.

## 3.4 Deleting a task

- A task is automatically deleted when the last line of its task code is encountered
- One can force a task to be deleted with *OSTaskDel* (*INT8U prio*) where *prio* is the priority of the task to delete.

## 3.4 Deleting a task

### *Warning*

- Sometimes, a task owns resources such as memory buffers or a semaphore.
- If another task attempts to delete this task, the resource will stay used (locked). Therefore, resource is lost.
- A good practice is to use *OSTaskDelReq(INT8U prio)*:
  - If Task A wants to delete Task B using *OSTaskDelReq*, then A asks B to *release all resources first and then to delete itself after*.

## 3.4 Deleting a task

### *Case of VxWorks*

- Before the acquisition of a resource (semaphore, mutex, etc.), the system function *TaskSafe()* can be used in order to avoid a deletion from other tasks.

## 3.4 Deleting a task

### *Another good practice*

- **Allocation and deallocation of TCB and stack are managed by the  $\mu$ C/OS-II**
- **But a task is responsible to:**
  - deallocate its own allocated memory (malloc, new)
  - properly release shared resource mechanisms (mutex, semaphore, etc.)
  - close all files
  - ...

## 3.5 Suspending/resuming a task

- Sometimes, it is useful to suspend the execution of a task explicitly. The *Suspended State* is included within the *Waiting State*.
- A suspended task can be resumed later
- With  $\mu\text{C}/\text{OS-II}$ :
  - *OSTaskSuspend(INT8Uprio)*
  - *OSTaskResume (INT8U prio)*



## 3.6 Delaying a task

- A task must have a mechanism to delay itself for an integral number of timer ticks.
- Very useful for a periodic execution or a busy waiting (read the status of an external device periodically).
- *Delayed state* is included in the *waiting state*.
- In  $\mu\text{C}/\text{OS-II}$ , for a task to delay itself:
  - *OSTimeDly(ticks)*, where *ticks* is the number of clock ticks from the timer;
  - *OSTimeDlyHMSM(H,M,S,M)*, with *H*ours, *M*inutes, *S*econds, *M*icroseconds

## 3.7 Changing the priority of a task

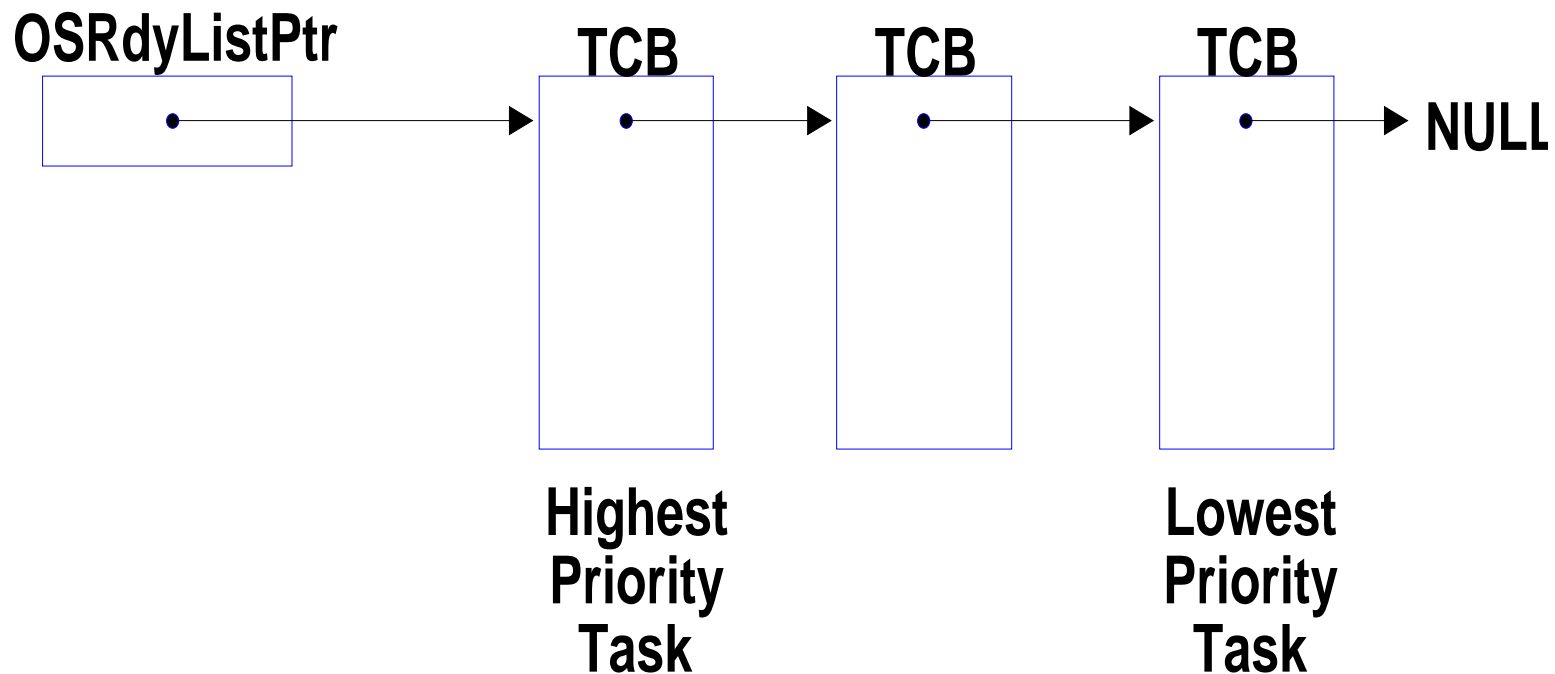
- When you create a task, you assign the task priority. At runtime, you can change the priority of any task.
- In other words, you can change priorities dynamically
- In  $\mu\text{C}/\text{OS-II}$ , for a task to change its priority:
  - *OSTaskChange(oldprio, newprio);*

## 3.8 Tasks scheduling

- A queue is used to maintain tasks that are ready to run (similarly, waiting tasks are memorized in queues, see chapter 5).
- The queue of “ready to run” tasks is sorted by priorities. That is, the highest priority task ready to run is selected first.
- More precisely, the type of queue is a bitmap. It allows the removing/deleting in  $O(\text{constant})$ .

## 3.8 Tasks scheduling

### *Queue of tasks ready to run*

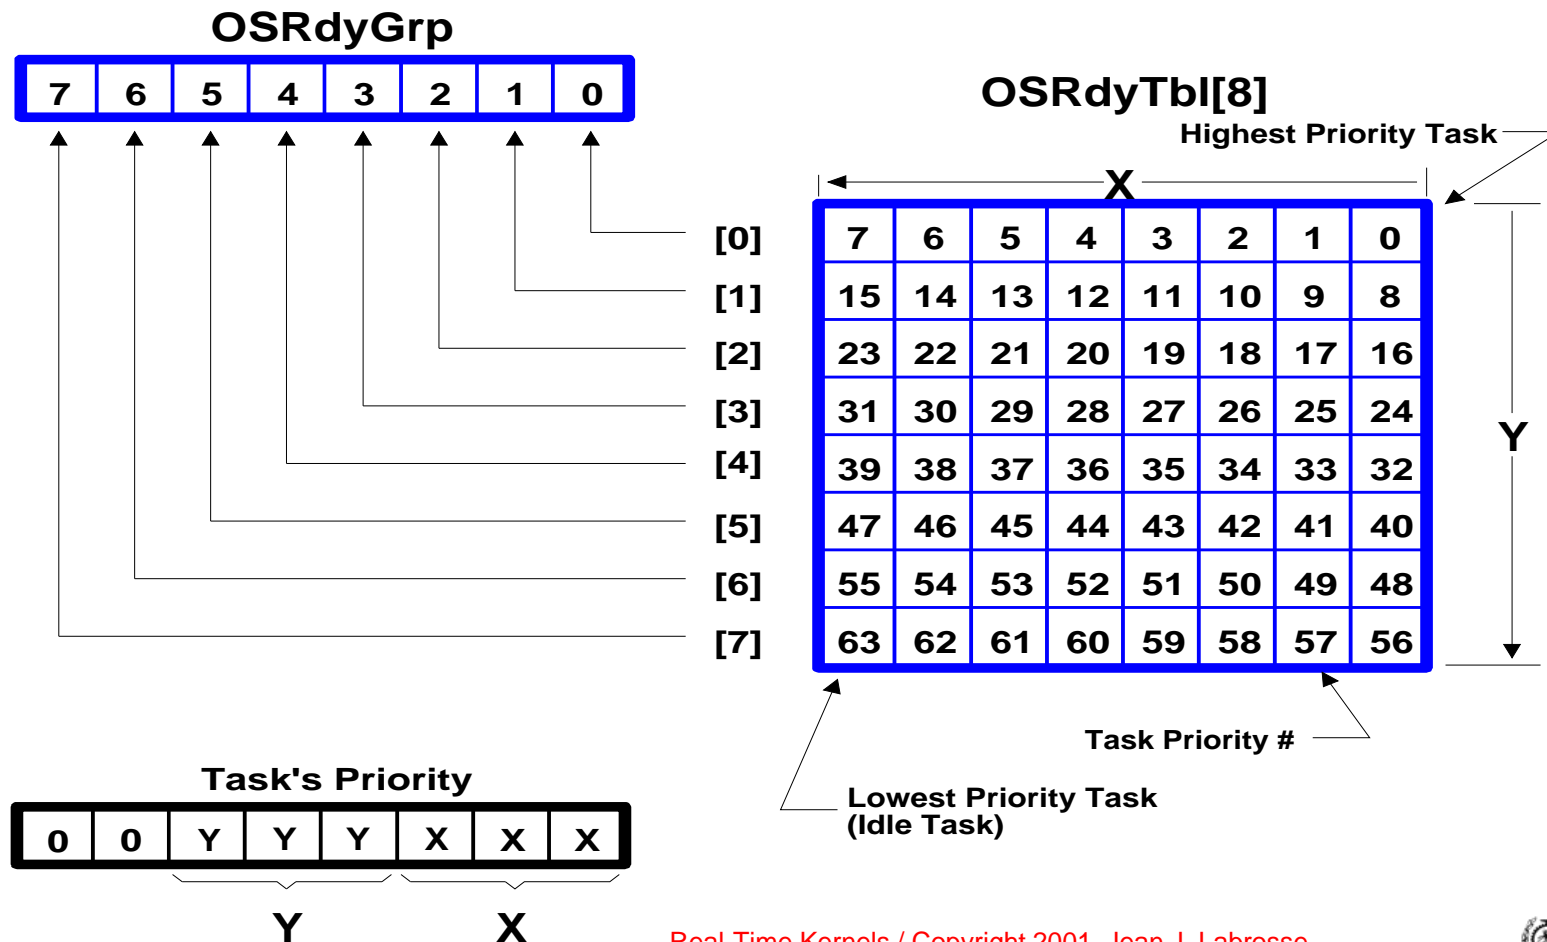


**(Running Task)**

Real-Time Kernels / Copyright 2001, Jean J. Labrosse

## 3.8 Tasks scheduling

### *Organization of priorities*





Real-Time Kernels / Copyright 2001, Jean J. Labrosse

Copyright© 2011 G. Bois, M. De Nanclas, L. Filion

## 3.8 Tasks scheduling

### Example



Task's Priority = ( Y \* 8 ) + X  
 Task's Priority = ( 1 \* 8 ) + 3  
 Task's Priority = 11

## 3.8 Tasks scheduling

### *Lookup table to resolve highest priority*

\*\*\*\*\*

\* Note(s): 1) Index into table is bit pattern to resolve highest priority.

\* 2) Indexed value corresponds to highest priority bit position

\*(i.e. 0..7)

\*\*\*\*\*

```

INT8U const OSUnMapTbl[] = {
    0, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x00-0x0F
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x10-0x1F
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x20-0x2F
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x30-0x3F
    6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x40-0x4F
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x50-0x5F
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x60-0x6F
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x70-0x7F
    7, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x80-0x8F
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0x90-0x9F
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0xA0-0xAF
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0xB0-0xBF
    6, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0xC0-0xCF
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0xD0-0xDF
    5, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0, // 0xE0-0xEF
    4, 0, 1, 0, 2, 0, 1, 0, 3, 0, 1, 0, 2, 0, 1, 0 // 0xF0-0xFF }

```

**X = @ [0x78]**  
(i.e. 0x78 = OSRdyTbl[1])

**Y = @ [0xF6]**  
(i.e. 0xF6 = OSRdyGrp)

## 3.8 Tasks scheduling

*Code to find the task ready to run with the highest priority*

```
Y          = OSUnMapTbl[OSRdyGrp];  
X          = OSUnMapTbl[OSRdyTbl[Y]];  
HighestPriority = (Y * 8) + X;
```

```
Y (i.e. 1)      = OSUnMapTbl[0xF6];  
X (i.e. 3)      = OSUnMapTbl[0x78];  
HighestPriority = (1 * 8) + 3;  
HighestPriority = 11
```



## 3.8 Tasks scheduling

- 4 constants (memorized in the TCB) are used to speedup access to OSRdyTbl and OSRdyGr:

OSTCBY  $\Rightarrow$  `ptcb->OSTCBY = priority >> 3;`

OSTCBBitY  $\Rightarrow$  `ptcb->OSTCBBitY = OSMaTbl[ptcb->OSTCBY]`

OSTCBX  $\Rightarrow$  `ptcb->OSTCBX = priority & 0x07;`

OSTCBBitX  $\Rightarrow$  `ptcb->OSTCBBitX = OSMaTbl[ptcb->OSTCBX];`

Index	Bit Mask
0	00000001
1	00000010
2	00000100
3	00001000
4	00010000
5	00100000
6	01000000
7	10000000

OSMaTbl[]

## 3.8 Tasks scheduling

- **Code to remove a task from the ready list:**

```
if ((OSRdyTbl[ptcb->OSTCBY] &= ~ ptcb->OSTCBBitX) ==0)
{
    OSRdyGrp &= ~ ptcb->OSTCBBitY;
}
```

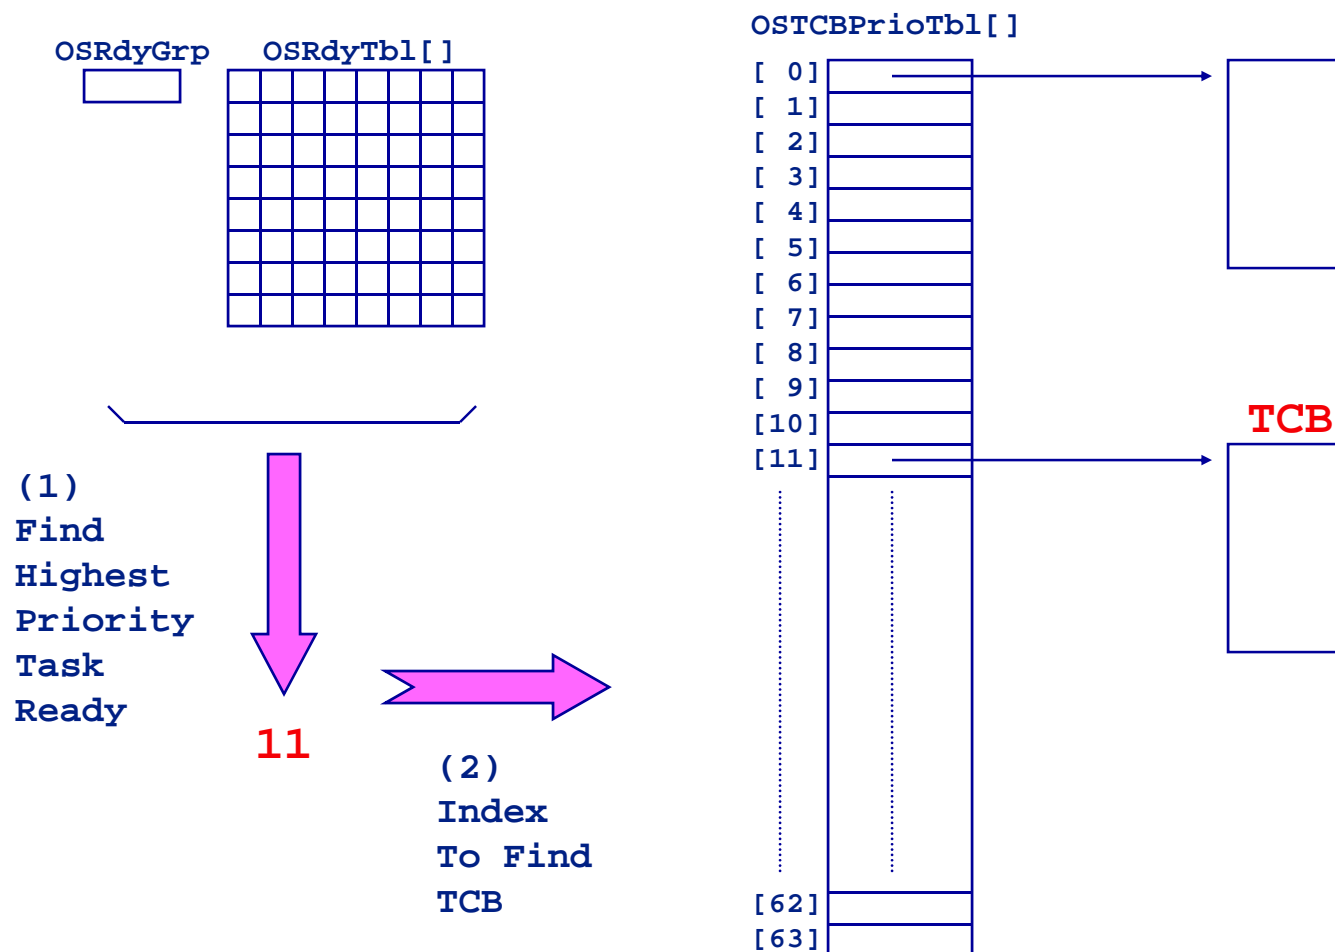
## 3.8 Tasks scheduling

- **Code to make a task ready to run**

```
OSRdyGrp          |=      ptcb->OSTCBBitY;  
OSRdyTbl[ptcb->OSTCBY] |=  ptcb->OSTCBBitX;
```

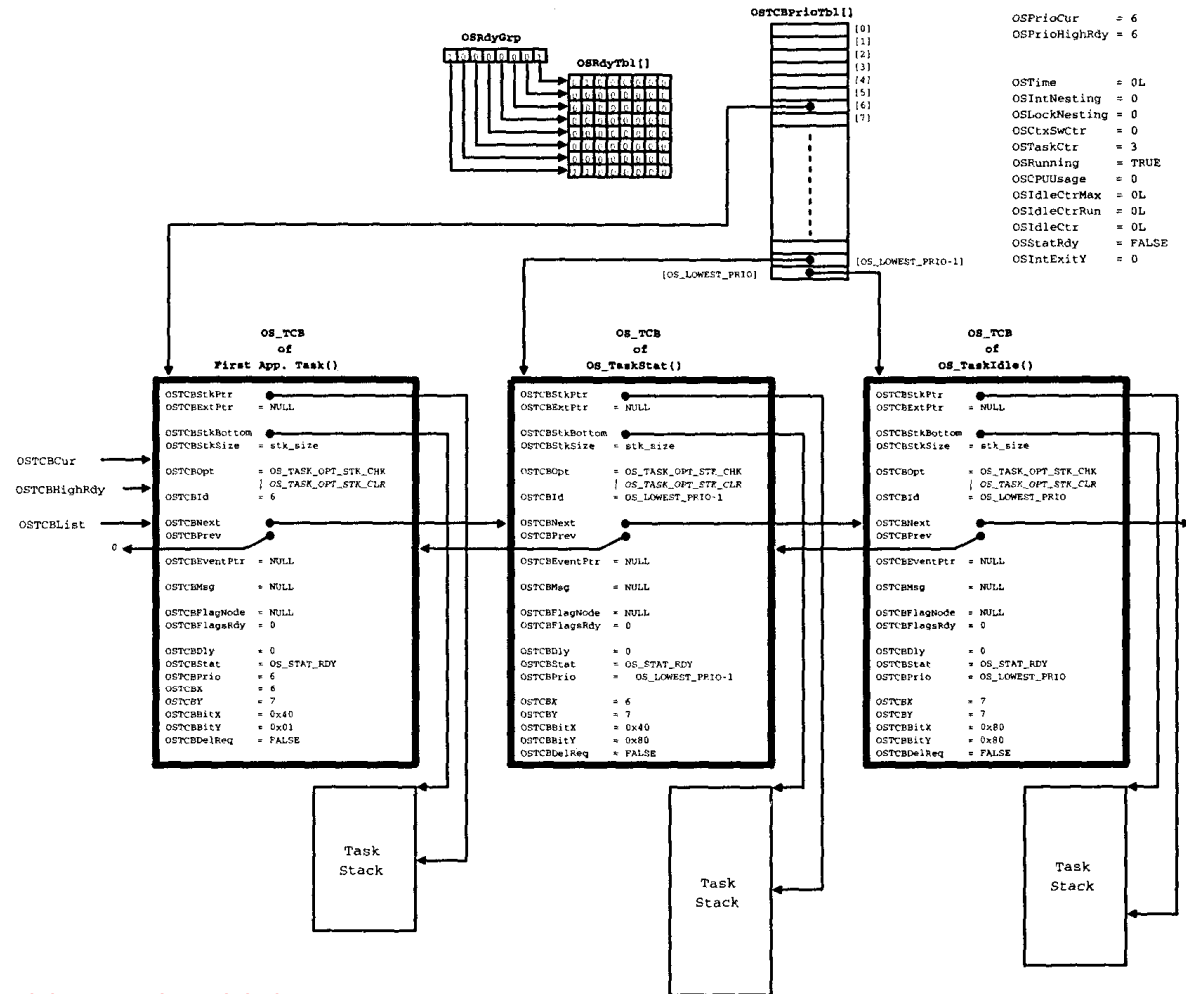
## 3.8 Tasks scheduling

*Highest priority task ready to run to its TCB*



# In more details

Variables and data structures after calling *OSStart()*.



Real-Time Kernels / Copyright 2001, Jean J. Labrosse

## 3.9 Systems tasks

- **RTOS includes several system tasks:**
  - Tasks used to log system messages without having to perform I/O in the current task context. (ex. VxWorks: tLogTask)
  - Task to execute function that cannot occur at interrupt level. It must have the highest priority in the system (ex. VxWorks: tExcTask)
  - Task executed when none other are ready to run. (*idle task*) must be the lowest priority
  - Tasks to provide run-time statistics (*statistics task*)

## 3.9 Systems tasks

### *Idle task*

- The lowest priority in  $\mu\text{C}/\text{OS-II}$  is *OS\_LOWEST\_PRIO* (e.g. 63).
- The task named *OS\_TaskIdle()* can never be deleted by applications software.
- *OS\_TaskIdle()* increments a 32-bit counter called *OSIdleCtr*.

## 3.9 Systems tasks

### *Statistics task*

- In  $\mu$ C this task is called *OS\_TaskStat()* and is created if you set the configuration constant *OS\_TASK\_STAT\_EN* to 1.
- When enabled, *OS\_TaskStat()* executes every second and computes the percentage of CPU usage.
- The priority of the statistics task is *OS\_LOWEST\_PRIO-1*.



## 3.9 Systems tasks

### *Statistics task*

- To compute the percentage of CPU usage, you must initially call the function *OSStatInit()* from the highest priority task.
- *OSStatInit()* determines the maximum value that a 32-bit counter (named *OSIdleCtrMax*) can reach during 1 second when no other task is executed.

## 3.9 Systems tasks

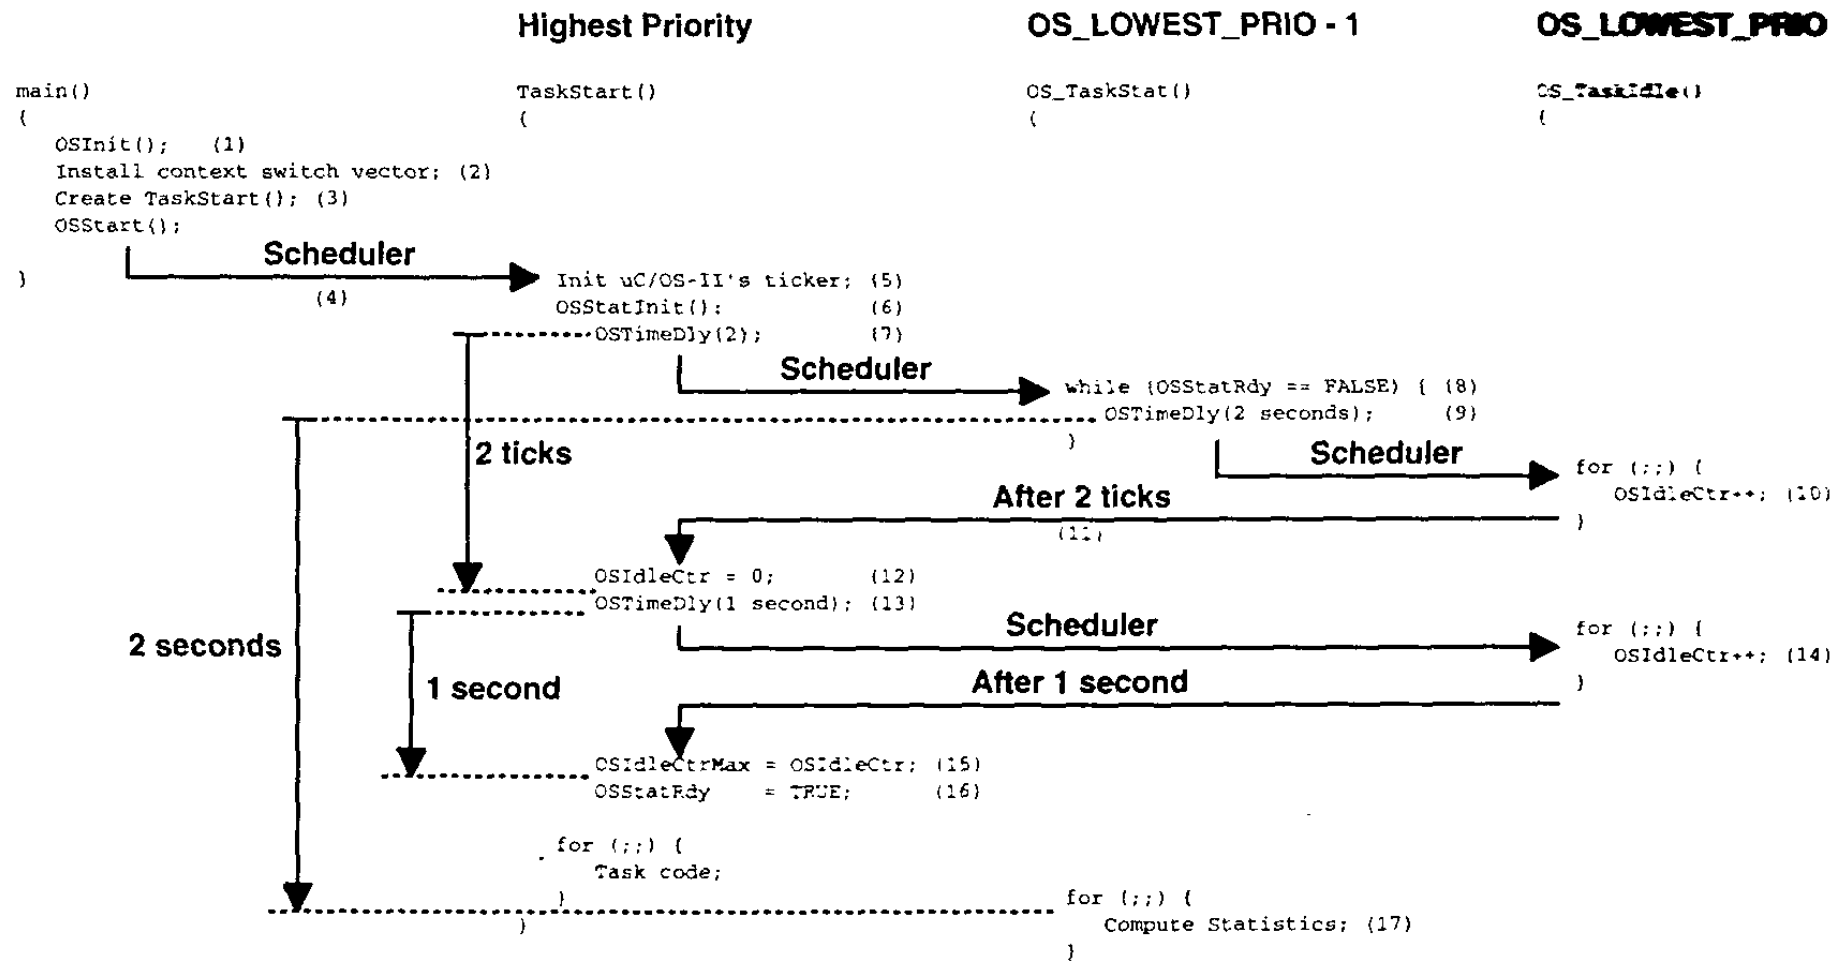
### *Statistics task*

- Each time the statistics task is executed, *OS\_TaskStat()* compares *OSIdleCtrMax* and *OSIdleCtr*.

$$OSCPUUsage(\%) = 100 * (1 - OSIdleCtr / OSIdleCtrMax)$$

- That is, CPU usage is stored in the variable *OSCPUUsage*

## Statistic task initialization.



Real-Time Kernels / Copyright 2001, Jean J. Labrosse